

Task: Perform EDA with the help of PowerBI

Task description:

XYZ is a private company in US which would to make investments in Cab industry, but before that they want to understand the market. In case of a positive decision, XYZ has to choos between two cab companies - 'Pink' and 'Yellow'.

Data description:

There are 4 data sets that contains information on 2 cab companies:

- 'Cab_Data.csv' - details of transactions
- 'Customer_ID.csv' - mapping table that contains a unique identifier which links the customer's demographic details
- 'Transaction_ID.csv' - mapping table that contains transaction to customer mapping and payment mode
- 'City.csv' - list of US cities, their population and number of cab users

Each file (data set) represents different aspects of the customer profile.

Preparation steps:

1. Import libraries:

```
In [1]: import numpy as np
import pandas as pd

import os
from datetime import datetime, timedelta

#!/pip install powerbiclient
from powerbiclient import Report
from io import StringIO
from ipywidgets import interact
```

2. Load data:

```
In [2]: cab_data = pd.read_csv('Cab_Data.csv', parse_dates=['Date of Travel'])
customer_ID = pd.read_csv('Customer_ID.csv')
transaction_ID = pd.read_csv('Transaction_ID.csv')
city = pd.read_csv('City.csv')
```

3. Explore tables

In [3]: cab_data

Out[3]:

	Transaction ID	Date of Travel	Company	City	KM Travelled	Price Charged	Cost of Trip
0	10000011	42377	Pink Cab	ATLANTA GA	30.45	370.95	313.6350
1	10000012	42375	Pink Cab	ATLANTA GA	28.62	358.52	334.8540
2	10000013	42371	Pink Cab	ATLANTA GA	9.04	125.20	97.6320
3	10000014	42376	Pink Cab	ATLANTA GA	33.17	377.40	351.6020
4	10000015	42372	Pink Cab	ATLANTA GA	8.73	114.62	97.7760
...
359387	10440101	43108	Yellow Cab	WASHINGTON DC	4.80	69.24	63.3600
359388	10440104	43104	Yellow Cab	WASHINGTON DC	8.40	113.75	106.8480
359389	10440105	43105	Yellow Cab	WASHINGTON DC	27.75	437.07	349.6500
359390	10440106	43105	Yellow Cab	WASHINGTON DC	8.80	146.19	114.0480
359391	10440107	43102	Yellow Cab	WASHINGTON DC	12.76	191.58	177.6192

359392 rows × 7 columns

```
In [4]: customer_ID
```

Out[4]:

	Customer ID	Gender	Age	Income (USD/Month)
0	29290	Male	28	10813
1	27703	Male	27	9237
2	28712	Male	53	11242
3	28020	Male	23	23327
4	27182	Male	33	8536
...
49166	12490	Male	33	18713
49167	14971	Male	30	15346
49168	41414	Male	38	3960
49169	41677	Male	23	19454
49170	39761	Female	32	10128

49171 rows × 4 columns

```
In [5]: transaction_ID
```

Out[5]:

	Transaction ID	Customer ID	Payment_Mode
0	10000011	29290	Card
1	10000012	27703	Card
2	10000013	28712	Cash
3	10000014	28020	Cash
4	10000015	27182	Card
...
440093	10440104	53286	Cash
440094	10440105	52265	Cash
440095	10440106	52175	Card
440096	10440107	52917	Card
440097	10440108	51587	Card

440098 rows × 3 columns

In [6]: city

Out[6]:

	City	Population	Users
0	NEW YORK NY	8,405,837	302,149
1	CHICAGO IL	1,955,130	164,468
2	LOS ANGELES CA	1,595,037	144,132
3	MIAMI FL	1,339,155	17,675
4	SILICON VALLEY	1,177,609	27,247
5	ORANGE COUNTY	1,030,185	12,994
6	SAN DIEGO CA	959,307	69,995
7	PHOENIX AZ	943,999	6,133
8	DALLAS TX	942,908	22,157
9	ATLANTA GA	814,885	24,701
10	DENVER CO	754,233	12,421
11	AUSTIN TX	698,371	14,978
12	SEATTLE WA	671,238	25,063
13	TUCSON AZ	631,442	5,712
14	SAN FRANCISCO CA	629,591	213,609
15	SACRAMENTO CA	545,776	7,044
16	PITTSBURGH PA	542,085	3,643
17	WASHINGTON DC	418,859	127,001
18	NASHVILLE TN	327,225	9,270
19	BOSTON MA	248,968	80,021

Comment: It is obvious, that tables *cab_data*, *transaction_ID* and *customer_ID* should be merged for the further analysis:

- column 'Transaction ID' joins *cab_data* and *transaction_ID*
- column 'Customer ID' joins *transaction_ID* and *customer_ID*

4. Merging tables:

```
In [7]: merged_data = pd.merge(cab_data, transaction_ID, on = 'Transaction ID', how = 'inner')
merged_data = pd.merge(merged_data, customer_ID, on = 'Customer ID', how = 'inner')
```

Cleaning data:

1. Check data types

```
In [8]: merged_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 359392 entries, 0 to 359391
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Transaction ID         359392 non-null  int64
1   Date of Travel         359392 non-null  object
2   Company               359392 non-null  object
3   City                  359392 non-null  object
4   KM Travelled           359392 non-null  float64
5   Price Charged          359392 non-null  float64
6   Cost of Trip           359392 non-null  float64
7   Customer ID            359392 non-null  int64
8   Payment_Mode           359392 non-null  object
9   Gender                 359392 non-null  object
10  Age                    359392 non-null  int64
11  Income (USD/Month)     359392 non-null  int64
dtypes: float64(3), int64(4), object(5)
memory usage: 35.6+ MB
```

All data types are correct except of the *Date of Travel*. This column still reflect the quantity of days from 1900 till today.

Here is how to change it:

```
In [9]: date = pd.to_datetime('1900-01-01') #count from this date
merged_data['Date of Travel'] = pd.to_numeric(merged_data['Date of Travel']) #covert column to numeric
merged_data['Date of Travel'] = merged_data['Date of Travel'].apply(lambda x: date + timedelta(x)) #calculate dates
```

In [10]: merged_data

Out[10]:

	Transaction ID	Date of Travel	Company	City	KM Travelled	Price Charged	Cost of Trip	Customer ID	Payment_Mode	Gender	Age	Income (USD/Month)
0	10000011	2016-01-10	Pink Cab	ATLANTA GA	30.45	370.95	313.6350	29290	Card	Male	28	10813
1	10351127	2018-07-23	Yellow Cab	ATLANTA GA	26.19	598.70	317.4228	29290	Cash	Male	28	10813
2	10412921	2018-11-25	Yellow Cab	ATLANTA GA	42.55	792.05	597.4020	29290	Card	Male	28	10813
3	10000012	2016-01-08	Pink Cab	ATLANTA GA	28.62	358.52	334.8540	27703	Card	Male	27	9237
4	10320494	2018-04-23	Yellow Cab	ATLANTA GA	36.38	721.10	467.1192	27703	Card	Male	27	9237
...
359387	10439790	2018-01-09	Yellow Cab	SEATTLE WA	16.66	261.18	213.9144	38520	Card	Female	42	19417
359388	10439799	2018-01-05	Yellow Cab	SILICON VALLEY	13.72	277.97	172.8720	12490	Cash	Male	33	18713
359389	10439838	2018-01-06	Yellow Cab	TUCSON AZ	19.00	303.77	232.5600	41414	Card	Male	38	3960
359390	10439840	2018-01-08	Yellow Cab	TUCSON AZ	5.60	92.42	70.5600	41677	Cash	Male	23	19454
359391	10439846	2018-01-06	Yellow Cab	TUCSON AZ	13.30	244.65	180.3480	39761	Card	Female	32	10128

359392 rows × 12 columns

2. Check and remove duplicates


```
In [11]: duplicates_indecies = np.where(merged_data.duplicated() == True)[0]

if len(duplicates_indecies) == 0:
    print("There is no duplicates")
else:
    print("Indecies:", duplicates_indecies)
```

There is no duplicates

3. Check and remove N/A

```
In [12]: bool_na = merged_data.isnull().values.any()

if bool_na == True:
    print("There are missed values in dataset")
else:
    print("There are no missed values")
```

There are no missed values

Exploratory Data Analysis with Power BI

Download merged data as .csv to visualize in Power BI:

```
In [13]: import os
os.makedirs('data/', exist_ok = True)
merged_data.to_csv('data/merged.csv')
city.to_csv('data/city.csv')
```

Setting parametrs for PowerBI:

```
In [14]: # Import the DeviceCodeLoginAuthentication class to authenticate against Power BI
from powerbiclient.authentication import DeviceCodeLoginAuthentication

# Initiate device authentication
device_auth = DeviceCodeLoginAuthentication()
access_token = device_auth.get_access_token()
```

Performing interactive authentication. Please follow the instructions on the terminal.

To sign in, use a web browser to open the page <https://microsoft.com/devicelogin> (<https://microsoft.com/devicelogin>) and enter the code DL7GW7SYJ to authenticate.

You have logged in.

Interactive authentication successfully completed.

Report data:

```
In [15]: group_id = "30c9ae19-4d19-4b04-b186-602205d30f72"
report_id = "eb2d3182-2996-4593-ab85-941ffc05a76b"
```

Load report:

```
In [16]: report = Report(group_id=group_id, report_id=report_id, access_token = access_token)
```

Click on visuals and walk through tabs.

Comments are in presentation.

- Simple statistics

In [17]: report

Simple Statistics

Date of Travel

1/4/2016



12/31/2018



46.13K

Quantity of Unique Customers

13

Quantity of States

18

Quantity of Cities

8.10M

Total KM Travelled

NEW YORK

brought biggest revenue

54

Max travels by one customer

NEW YORK

highest revenue per km

151.99M

Total USD charged

102.71M

Total USD spent

49.28M

Total Revenue


15.05K

Average of Income (USD/Month)

Geography insights


State	City	Unique customers
NY	NEW YORK	99744
IL	CHICAGO	56544
CA	LOS ANGELES	47966
DC	WASHINGTON	43673
MA	BOSTON	29659
CA	SAN DIEGO	20461
CA	SILICON VALLEY	12482
WA	SEATTLE	7986
Total		358879

Average of Price Charged, Average of Cost of Trip and Average of Revenue by State



Total Revenue by Year, Month and State

State: AZ CA CO DC FL GA IL MA NY PA TN TX WA



49.28M


Total Revenue

8.10M


KM Travelled

Sociological data


Quantity of Unique Customers by Gender




Customers by Income group (average value for group)



Quantity of Unique Customers by Payment Mode



Quantity of customers by Age and Gender



35.34

Average of Age

15.05K


Average of Income (USD/Month)

In [20]: report

Data by Company


Unique customers by Year, Month and Company

Company ● Pink Cab ● Yellow Cab



Average of Revenue per km by Year, Month and Company

Company ● Pink Cab ● Yellow Cab



8.10M

KM Travelled

46.13K

Unique customers


Company

□ Pink Cab

□ Yellow Cab

Cost of km and Revenue per km by Year and Month

● Cost of km ● Revenue per km



19.00

Min of Cost of Trip

691.20

Max of Cost of Trip